

PiotrKuczko

March 11, 2021

1 House Sales in King County, USA

1.0.1 Predict house price using regression

```
[21]: import numpy as np
import pandas as pd
import xgboost as xg
from sklearn.model_selection import train_test_split
from sklearn.metrics import mean_squared_error as MSE
from sklearn import preprocessing
import xgboost as xgb
from xgboost.sklearn import XGBRegressor
import datetime
from sklearn.model_selection import GridSearchCV
```

```
[22]: df = pd.read_csv('kc_house_data.csv')
df
```

```
[22]:
```

	id	date	price	bedrooms	bathrooms	\
0	7129300520	20141013T000000	221900.0	3	1.00	
1	6414100192	20141209T000000	538000.0	3	2.25	
2	5631500400	20150225T000000	180000.0	2	1.00	
3	2487200875	20141209T000000	604000.0	4	3.00	
4	1954400510	20150218T000000	510000.0	3	2.00	
...	
21608	263000018	20140521T000000	360000.0	3	2.50	
21609	6600060120	20150223T000000	400000.0	4	2.50	
21610	1523300141	20140623T000000	402101.0	2	0.75	
21611	291310100	20150116T000000	400000.0	3	2.50	
21612	1523300157	20141015T000000	325000.0	2	0.75	

	sqft_living	sqft_lot	floors	waterfront	view	...	grade	\
0	1180	5650	1.0	0	0	...	7	
1	2570	7242	2.0	0	0	...	7	
2	770	10000	1.0	0	0	...	6	
3	1960	5000	1.0	0	0	...	7	
4	1680	8080	1.0	0	0	...	8	
...	

21608	1530	1131	3.0	0	0	...	8
21609	2310	5813	2.0	0	0	...	8
21610	1020	1350	2.0	0	0	...	7
21611	1600	2388	2.0	0	0	...	8
21612	1020	1076	2.0	0	0	...	7

	sqft_above	sqft_basement	yr_built	yr_renovated	zipcode	lat	\
0	1180	0	1955	0	98178	47.5112	
1	2170	400	1951	1991	98125	47.7210	
2	770	0	1933	0	98028	47.7379	
3	1050	910	1965	0	98136	47.5208	
4	1680	0	1987	0	98074	47.6168	
...	
21608	1530	0	2009	0	98103	47.6993	
21609	2310	0	2014	0	98146	47.5107	
21610	1020	0	2009	0	98144	47.5944	
21611	1600	0	2004	0	98027	47.5345	
21612	1020	0	2008	0	98144	47.5941	

	long	sqft_living15	sqft_lot15
0	-122.257	1340	5650
1	-122.319	1690	7639
2	-122.233	2720	8062
3	-122.393	1360	5000
4	-122.045	1800	7503
...
21608	-122.346	1530	1509
21609	-122.362	1830	7200
21610	-122.299	1020	2007
21611	-122.069	1410	1287
21612	-122.299	1020	1357

[21613 rows x 21 columns]

```
[23]: df[['year', 'month', 'day']] = pd.DataFrame([ [int(x[0:4]), int(x[4:6]),
→int(x[6:8])] for x in df['date'].tolist() ])
df = df.drop(['date', 'id'], axis=1)
Y = df['price']
X = df.drop(['price'], axis=1)
```

```
[24]: train_X, test_X, train_Y, test_Y = train_test_split(X, Y,
test_size = 0.3, random_state = 123)
```

```
[28]: xgb1 = XGBRegressor()
parameters = {'nthread':[1], #when use hyperthread, xgboost may become slower
'objective':['reg:squarederror'],
'learning_rate': [.03, .04, .05, .06, .07], #so called `eta` value
```

```

        'max_depth': [3, 4, 5, 6, 7],
        'min_child_weight': [3, 4, 5, 6],
        'subsample': [0.7],
        'colsample_bytree': [0.7],
        'n_estimators': [500, 700, 100]}

xgb_grid = GridSearchCV(xgb1,
                        parameters,
                        cv = 2,
                        n_jobs = 12,
                        verbose=True)

xgb_grid.fit(X, Y)

print(xgb_grid.best_score_)
print(xgb_grid.best_params_)

```

Fitting 2 folds for each of 300 candidates, totalling 600 fits
0.877865911661787
{'colsample_bytree': 0.7, 'learning_rate': 0.06, 'max_depth': 5,
'min_child_weight': 6, 'n_estimators': 700, 'nthread': 1, 'objective':
'reg:squarederror', 'subsample': 0.7}

```

[31]: xgb2 = XGBRegressor(colsample_bytree=0.7, learning_rate=0.06, max_depth=5,
    ↪min_child_weight=6, n_estimators=700, nthread=1, objective='reg:
    ↪squarederror', subsample=0.7)

```

```

[36]: xgb2.fit(train_X, train_Y)
score = xgb2.score(train_X, train_Y)
print("Training score: ", score)
score = xgb2.score(test_X, test_Y)
print("Test score: ", score)
pred = xgb2.predict(test_X)
rmse = np.sqrt(MSE(test_Y, pred))
print("RMSE : % f" %(rmse))

```

Training score: 0.9711698297228819
Test score: 0.905722369519171
RMSE : 113149.749407

```

[37]: print (test_Y[:10], pred[:10])

```

```

5506      532500.0
9279      410000.0
16034     782500.0
6608      995000.0
20359     279000.0
8798      175000.0

```

```


10035    689000.0
13321    275000.0
15842    465000.0
12119    506000.0
Name: price, dtype: float64 [617686.7  525400.25 869056.5  986680.4  302756.4
301077.2  696768.8
419573.34 457880.12 462333.66]

```

```

[45]: import matplotlib.pyplot as pyplot
pyplot.bar(X.columns, xgb2.feature_importances_)
pyplot.xticks(rotation=90)
pyplot.show()

```



```
[ ]:
```